id: letsvibeai-learning-workspace-planner
name: LetsVibeAI Learning Workspace Planner
version: 1.0.0
status: draft
owner: rostr-pal-skill-builder
category: planning
trigger: A LetsVibeAI product, growth, curriculum, workspace, or sandbox decision needs a grounded plan and a testable next step.
inputs:

- name: product_request
  required: true
  description: The proposed product, growth, curriculum, workspace, or sandbox change.
- name: project_context
  required: true
  description: Current course, plans, known constraints, metrics, and approved technology boundaries.
  outputs:
- name: decision_package
  format: markdown
  description: A traceable recommendation with plan, risks, measures, and approval gates.
  allowed_tools:
- project-file-search
- approved-web-research
- artifact-store
- analytics-read
- sandbox-read
  denied_tools:
- production-deployment
- billing-changes
- external-messaging
- secret-access
- destructive-data-actions
  requires_approval_for:
- pricing-or-billing-change
- external-file-ingestion
- public-deployment
- third-party-account-connection

- production-sandbox-execution
  memory_namespace: project/{project_id}

LetsVibeAI Learning Workspace Planner

Purpose

Turn a LetsVibeAI opportunity into the smallest valuable, safe, revenue-aware learning-and-building experiment.

Use when

- Planning a course, lab, workspace, chat, sandbox, monetization, onboarding, retention, or project-template change.
- A learner request must become a PAL project brief and a structured build path.

Do not use when

- The request is a simple editorial correction with no impact on learning flow, product behavior, pricing, or safety.
- The action would directly modify production, bill a user, or publish externally; route that action to the appropriate approved execution workflow.

Inputs

Input	Required	Validation	Notes
product_request	Yes	States desired outcome or observed problem	Preserve original wording
project_context	Yes	Includes known constraints and existing artifacts	AWS-only unless explicitly changed
evidence	No	Links/metrics/source register	Required for external or current claims

Outputs

Output	Format	Definition of done
decision_package	Markdown	Includes scope, JTBD, NPAO priority, MVP, measures, risks, and next action
project_contract_update	YAML	Labels stated requirements, assumptions, approvals, and owners

Procedure

1. Classify the request as acquisition, activation, learning, monetization, retention, sandbox safety, or platform capability.
2. Extract stated requirements, inferred requirements, non-goals, constraints, and material open decisions.
3. Write user, system, build, and operational JTBD statements.

4. Check available product analytics and relevant course/project artifacts before external research.
5. If evidence is needed, create tiered RAG-DAL research targets; do not treat community commentary as sole support for a material decision.
6. Rank options with NPAO. Dependencies, security, learner trust, and explicit commitments override a numeric score.
7. Recommend the smallest reversible experiment with success metric, failure signal, owner, cost/usage boundary, and rollback.
8. Create approval requests for pricing, billing, public deployment, third-party connections, private data use, or consequential sandbox execution.
9. Update the project contract and return a concise decision package.

Guardrails

- Do not assume a paywall is justified; connect every paid feature to an ongoing value loop or measurable cost.
- Keep foundational education accessible; charge for persistent workspace, guidance, compute, templates, collaboration, or advanced execution value.
- Do not expose credentials in prompt, artifact, or log output. Use `ENV:<NAME>` references.
- Sandbox recommendations must use isolation, bounded resources, restricted egress, checkpointing, and teardown.
- Do not represent a recommendation, price, or projection as validated without evidence.

Quality checks

- [ ] The outcome and target learner are explicit.
- [ ] Every proposed feature maps to a job, metric, and test.
- [ ] The plan distinguishes existing facts, assumptions, and optional expansion.
- [ ] Cost-bearing resources have a quota/metering or approval policy.
- [ ] External or irreversible actions have a named approval gate.

Failure and escalation

Condition	Action
No baseline metrics	Define instrumentation first; do not claim conversion impact
Material audience ambiguity	Ask one focused learner/segment question
Sandbox safety requirement unclear	Restrict to a non-production template and request architecture review
Pricing decision required	Frame an experiment, not a permanent change; request approval before implementation

Examples

Input

```
product_request: Add a personal AI workspace so students can build apps while taking
project_context:
  curriculum: six free modules and three labs
  constraints: AWS-only; no production secrets in learner workspace
```

Expected output

```
recommendation: Launch a quota-bound Builder Workspace as a paid experiment after a 
now:
  - Implement project brief, workspace entitlement, one isolated starter sandbox, and
metric: project_created_to_first_sandbox_run
approval_required:
  - billing configuration
  - public sandbox deployment
```

Change log

- 1.0.0 — Initial LetsVibeAI integration skill.